KEENAN S. KEELING

Hyattsville, MD (Washington, DC area) · 301-461-2163 · kkeeling@gmail.com · linkedin.com/in/keenankeeling · github.com/kkeeling

Engineering leader · 25+ years shipping production software · Multiplies engineering teams with AI, under written controls

SUMMARY

Engineering leader with 25+ years shipping production software and three Director of Engineering roles across aviation, IoT, and enterprise AI. At Parslee.AI I built the system that took one team's merged output from 59 pull requests a month to 505, with zero reverts, by giving engineers AI agent fleets that run under an independent reviewer, mandatory test, lint, and supply-chain gates, and human release control. I do not replace engineers. I give them the tooling, the guardrails, and the habits to deliver at a multiple of what they ship today, and I can install that system in an engineering organization in its first quarter.

WHAT I BRING TO AN ENGINEERING ORGANIZATION

- **A proven multiplier for the same team.** Merged output grew from 59 pull requests in September 2025 to 505 in the 30 days to October 9, 2026, with 83% of that work produced by agent fleets that engineers direct and work landing on 29 of 30 days. Engineers set the goals and acceptance checks; the fleet carries the typing. [KEENAN: the essay says 473 PRs in 31 days (Aug 27 to Sep 26); reconcile the two windows so both documents agree.]
- **Definition of done before a line of code.** For anything larger than a routine change, agents write the outcomes and one yes/no check per outcome first. Example, September 18: five writer agents and a critic produced 235 outcomes and 215 checks for CAR's pull-request reviewer; 211 checks were met with evidence, six pull requests merged and deployed, and a cross-vendor verifier found two critical security holes that no same-vendor check caught.
- **Speed under enterprise controls.** Every agent-written change is reviewed by a separate reviewer agent that never reviews its own work, re-runs the evidence itself, and must pass a bench plus automated test, lint, and supply-chain checks before a guarded merge script will accept it. Anything a user sees is walked in a real browser before merge and again in production. In the last 30 days, 466 changes merged into the two main codebases with zero reverts and [N] production incidents.
- **Written authority, not blanket autonomy.** Every consequential action is sorted into ask-first (customer auth and billing, production credentials, permissions, releases, spend: 9 kinds), never (irreversible production deletes), do-then-tell (12 kinds), or just-do (8 kinds). Blockers reach a person as one line with options and a default; everything else keeps moving.
- **A system the whole team owns, not a hero.** Wheelhouse, a public template, installs a fleet (lead, builder, reviewer, checker) into any repository. goalpool queues work to agents and lets teammates dispatch and monitor it remotely. Fleets run in five codebases and other engineers operate them without me. When an agent repeats a mistake, the fix goes into shared rules and checks, so the lesson compounds.
- **Cost that survives a procurement review.** The fleets run on about $1,800 a month in flat-rate subscriptions; at API list pricing the same token volume would cost roughly $[X]K a month [KEENAN: Grace to compute from actual token usage], against the salary cost of the 25+ engineers the output would otherwise require.
- **Engineering judgment that scales.** Corrections go into each fleet's standing brief, dated, in the engineer's exact words, so the review comments senior engineers used to leave by hand now live in tests, lint rules, and shared contracts, and the same mistake stops coming back.

PROFESSIONAL EXPERIENCE

Parslee.AI — Director of Engineering

Remote · Sep 2025 – Present

Parslee puts AI employees inside a company's Microsoft 365, each with a role, a level of authority, and a memory. Policy-gated agents acting on email, calendar, calls, CRM, and documents, the same trust problem a records or evidence platform has.

- Grew the team's merged output from 59 to 505 pull requests a month (3,100+ merged since September 2025) by building and rolling out agent-fleet delivery across five codebases, with zero reverts in the two main codebases over the last 30 days.
- Built Wheelhouse, the public agent-fleet template (github.com/morebetterltd/wheelhouse) that turns a goal and its acceptance checks into builder, reviewer, and checker work, adopted by [N] engineers on the team.

- Contribute to CAR, the Common Agent Runtime that checks every proposed model action against policy before it runs, the control layer that lets AI act on customer data in a regulated enterprise. Built goalpool to queue and fail over agent work across accounts.
- Lead the engineering team [N engineers]: hiring, 1:1s, weekly priority-setting, and onboarding engineers onto fleet-based delivery [KEENAN: days to first production merge for the newest hire].

flyExclusive, Inc. — Director of Engineering Remote · Sep 2024 – Sep 2025

- Took the pilot mobile app from unused to installed and used regularly by every pilot [KEENAN: pilot count, and adoption or usage metric] by rebuilding it with a team of [N] engineers.

Volato, Inc. — Director of Software Engineering Remote · 2022 – Sep 2024

- Built and led a software engineering team of [N] across customer, crew, and internal systems [KEENAN: hires made, promotions].
- Delivered Fly Vaunt, an "empty leg" flight marketplace, generating $1.3M+ revenue and nearly 78,000 users in year one.
- Shipped Volato Mobile (11,000 users in six months) and the Volato Crew App for pilot scheduling (20 pilots in beta).
- Launched Volato AI, a GPT-4o sales agent inside the sales team's tools [KEENAN: sales outcome, e.g. response time or leads handled].
- Kept customer-facing and internal systems at near-100% uptime with monitoring dashboards, admin tooling, and on-call discipline.

Agrify, Inc. — Director of Software Engineering Remote (Burlington, MA) · 2018 – 2022

- Led a team of [N] engineers building the Java IoT controllers that run Agrify's indoor farming units, maintaining 100% unit uptime.
- Helped build Agrify Insights, award-winning farm-operations software that reached 90% facility utilization, cut operating costs 30%, and ingests 1.5M+ data points per customer per year.

MoreBetter, Ltd. (Releaf) — Co-founder Washington, DC · 2015 – Present (part-time)

- Built Releaf, the top-rated cannabis journal on the App Store (4.7), to 100,000+ users and nearly 900,000 journal entries.
- Named inventor on US patent 11,253,195 B2. Co-authored 15 published studies with the University of New Mexico and others.

Automata Studios — Co-founder, CEO & Software Engineer Washington, DC · 2006 – 2019

- Grew a software and hardware studio to $10M+ revenue over 12 years delivering for IBM Watson, Frank Gehry, Victoria's Secret, Barnes & Noble, Adobe, and MTV. Architected and built the Kids' Reader on Barnes & Noble's NOOK.

WORK YOU CAN INSPECT

- Wheelhouse, agent-fleet template (public, MIT, comes with me): github.com/morebetterltd/wheelhouse
- How I build with agent fleets, a 12-minute written tour of the setup, the controls, and a seven-stage path any engineer can follow: claude.ai/artifact/Ut67onpaXX98VadmYvfEEe
- Neo, code-reasoning engine with persistent semantic memory, co-built: github.com/Parslee-ai/neo
- A live walk-through of the fleets at work, on request

SKILLS

AI engineering: agent fleets, Claude Code, Model Context Protocol, LLM integration, policy-gated agent runtimes

Languages and data: C#/.NET, Java, TypeScript/Node.js, Python/Django, Swift, PostgreSQL, SQL, Cosmos DB

Cloud: Azure, AWS, Microsoft 365/Graph · Front end and mobile: React/Next.js, React Native, iOS

Engineering practice: architecture reviews, CI/CD with mandatory test, lint, and supply-chain gates, high-availability services, mentoring

EDUCATION

University of Maryland, College Park: BS, Computer Science